ENTERPRISE MANAGER DATABASE REPLAY EXECUTIVE SUMMARY

sh_upgradetest_02

sh_upgradetest_02 finished with status COMPLETED. Replay DB time regressed by 347.1% versus capture. Divergence is LOW at 2.79% of calls. Overall verdict: bad.

FAIL Bottom line: bad

RISK RATING

Moderate

REPLAY DATABASE

DB26(23.26.1.0.0)

CAPTURE DATABASE

SDB2(19.3.0.0.0)

Executive Summary

This Oracle Database Replay validated an upgrade from version **19.3.0.0.0** to **23.26.1.0.0**. The replay completed successfully; however, Database Time significantly increased by **347.1%**, indicating degraded performance in the replay environment. The replay system operated with fewer CPU cores and slightly less physical memory than the capture environment, which likely impacted performance comparability. Replay divergence was low at **2.79%**, supporting a generally reliable replay despite some differences. Parsing overhead increased substantially, with parse time share rising from **19.33%** to **43.94%** and total parses growing from 54,663 to 78,271. These factors highlight areas warranting further examination.

Test Outcome

- **Status:** COMPLETED

- **Was the test valid?** Yes - Replay completed and the core comparability checks do not show a major fidelity failure.

Test Objective

- **Change being tested:** Version upgrade
- **Reason:** Capture ran on **19.3.0.0.0** and replay ran on **23.26.1.0.0**.

Performance Assessment

- **Overall:** Poor
- **DB Time:** Degraded
- **User Calls:** Similar
- **Divergence:** Good
- **Session completion:** Stable

Key Findings

- Replay completed successfully, validating database upgrade from **19.3.0.0.0** to **23.26.1.0.0**.
- Database Time increased sharply by **347.1%** (from 7,572 to 33,858 seconds), indicating slowed replay performance.
- CPU resources were reduced from 2 cores in capture to 1 core in replay, potentially affecting throughput.
- Physical memory dropped slightly (15.33 GB to 15.17 GB), which may contribute to performance differences.
- Replay divergence level was low (**2.79%**), confirming acceptable replay fidelity.

- Parsing overhead grew considerably; parse time ratio more than doubled and total parses increased by over **40%**.
- Top replay wait event was latch contention on the shared pool, responsible for **35.56%** of replay Database Time.

Likely Causes

- Replay environment had reduced CPU core count compared to capture, impacting concurrency and throughput.
- Slightly lower physical memory on replay system may affect caching and buffer performance.
- Parameter changes associated with the upgrade potentially altering execution plans and optimizer behavior.
- High parsing intensity and literal SQL usage contributing to additional overhead and delays.

Final Verdict

❌ Unsuccessful test result

Replay status and fidelity checks are sufficient for a directional conclusion.

Why this result was assigned:

- Replay completed and the core comparability checks do not show a major fidelity failure.
- Capture ran on **19.3.0.0.0** and replay ran on **23.26.1.0.0**.
- Performance assessment for DB Time: **Degraded**.
- Divergence assessment: **Good**.

Diagnostic Notes

Moderate

Replay divergence is present

Issue: LOW divergence with 2.79% of calls diverging.

Likely cause: The main divergence signal comes from SELECT row-count differences, which often indicates data drift, optimizer plan changes, or environmental differences between capture and replay.

Recommendation: Validate schema statistics, application data parity, initialization parameters, and replay setup before drawing strong performance conclusions.

High

Replay performance regressed

Issue: Database Time increased by 347.12% in replay.

Likely cause: Replay spent more total DB time than capture, indicating reduced throughput or higher contention.

Recommendation: Investigate the replay's top waits, SQL regressions, and hardware saturation before promoting the tested change.

Low

Replay shifted toward CPU-heavy execution

Issue: CPU time increased by 7.4% and now represents 2.56% of replay DB time.

Likely cause: The target system is spending a larger share of time on CPU, which can happen after I/O bottlenecks are reduced, when SQL plans change, or when parsing and metadata work increase.

Recommendation: Review top SQL, execution plans, and CPU headroom to ensure the improvement does not create a new CPU ceiling.

Moderate

Concurrency waits remain visible in replay

Issue: latch: shared pool accounts for 35.56% of replay DB time with 31017 waits.

Likely cause: Library cache and row cache contention usually point to parse pressure, shared pool churn, metadata contention, or repeated object invalidation.

Recommendation: Reduce hard parsing, check shared pool sizing, stabilize SQL reuse, and review dictionary-heavy operations during replay.

Evidence Snapshot

Metric	Value
Replay status	COMPLETED
Replay divergence	LOW (2.79%)
Database Time change	347.12%
CPU Time change	7.37%
Replay average active sessions	70.61
AWR top replay wait	latch: shared pool
Generated	3/31/2026, 8:39:55 PM

Prototype Verdict Criteria

This section explains the prototype-only scoring logic used to derive the Bottom Line and Risk Rating. It is intended as reviewer guidance and methodology, not as part of the replay performance findings themselves.

How Bottom Line Is Defined

This prototype uses a rule-based scoring approach, not an LLM, to assign the Bottom Line. Reviewers can use the table below to see the explicit threshold logic.

Outcome	Prototype Rule

PASS / good	Assigned when Database Time improves materially, using DB Time change <= -**20%**, divergence stays below **3%**, Resource Manager CPU throttling stays below **10%**, and replay CPU capacity is not lower than capture.
WARN / mixed	Assigned when the result is inconclusive or partially positive: for example DB Time improves by less than **20%**, DB Time regression is **10%** or less, or comparability is weakened by throttling, CPU mismatch, or non-zero divergence.
FAIL / bad	Assigned when Database Time regresses materially, using DB Time change > +**10%**, or when the replay shows clear evidence that performance is worse in the test environment.

How Risk Rating Is Defined

Risk Rating is a separate roll-up of diagnostic findings. Each finding is assigned a severity and converted into a score:

- **High finding** = 35 points
- **Moderate finding** = 20 points
- **Low finding** = 8 points

The total score is then mapped to the displayed Risk Rating:

Risk Rating	Prototype Rule
Low	Total finding severity score below 60. Usually means only a few low or moderate issues were detected.
Moderate	Total finding severity score from 60 to 84. Usually means multiple moderate issues or one high issue plus other concerns were detected.

High	Total finding severity score 85 or above. Usually means several high-impact issues were detected in the replay analysis.